



GPU Teaching Kit

Accelerated Computing



Module 10 – Parallel Computation Patterns (scan)


Lecture 10.2 - A Work-inefficient Scan Kernel

Objective

- To learn to write and analyze a high-performance scan kernel
 - Interleaved reduction trees
 - Thread index to data mapping
 - Barrier Synchronization
 - Work efficiency analysis

A Better Parallel Scan Algorithm


1. Read input from device global memory to shared memory
2. Iterate $\log(n)$ times; stride from 1 to $n-1$: double stride each iteration



- Active threads *stride* to $n-1$ (n -stride threads)
- Thread j adds elements j and j -*stride* from shared memory and writes result into element j in shared memory
- Requires barrier synchronization, once before read and once before write


A Better Parallel Scan Algorithm

1. Read input from device to shared memory
2. Iterate $\log(n)$ times; stride from 1 to $n-1$: double stride each iteration.




A Better Parallel Scan Algorithm

1. Read input from device to shared memory
2. Iterate $\log(n)$ times; stride from 1 to $n-1$: double stride each iteration
3. Write output from shared memory to device memory



Handling Dependencies

- During every iteration, each thread can overwrite the input of another thread
 - Barrier synchronization to ensure all inputs have been properly generated
 - All threads secure input operand that can be overwritten by another thread
 - Barrier synchronization is required to ensure that all threads have secured their inputs
 - All threads perform addition and write output



ITERATION = 1

STRIDE = 1

A Work-Inefficient Scan Kernel

```
__global__ void work_inefficient_scan_kernel(float *X, float *Y, int InputSize) {
    __shared__ float XY[SECTION_SIZE];
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < InputSize) {XY[threadIdx.x] = X[i];}
    // the code below performs iterative scan on XY
    for (unsigned int stride = 1; stride <= threadIdx.x; stride *= 2) {
        __syncthreads();
        float in1 = XY[threadIdx.x - stride];
        __syncthreads();
        XY[threadIdx.x] += in1;
    }
    __syncthreads();
    If (i < InputSize) {Y[i] = XY[threadIdx.x];}
}
```

Work Efficiency Considerations

- This Scan executes $\log(n)$ parallel iterations
 - The iterations do $(n-1)$, $(n-2)$, $(n-4)$, ..., $(n - n/2)$ adds each
 - Total adds: $n * \log(n) - (n-1) \rightarrow O(n * \log(n))$ work
- This scan algorithm is not work efficient
 - Sequential scan algorithm does n adds
 - A factor of $\log(n)$ can hurt: 10x for 1024 elements!
- A parallel algorithm can be slower than a sequential one when execution resources are saturated from low work efficiency



GPU Teaching Kit



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).