


ADVANCED OPENACC COURSE

Lecture 1: Advanced Profiling, May 19, 2016



Course Objective:

Enable *you* to scale *your* applications on multiple GPUs and optimize with profiler tools

Course Syllabus

May 19: Advanced Profiling of OpenACC Code

May 26: Office Hours

June 2: Advanced multi-GPU Programming with MPI and OpenACC

June 9: Office Hours

Recordings:

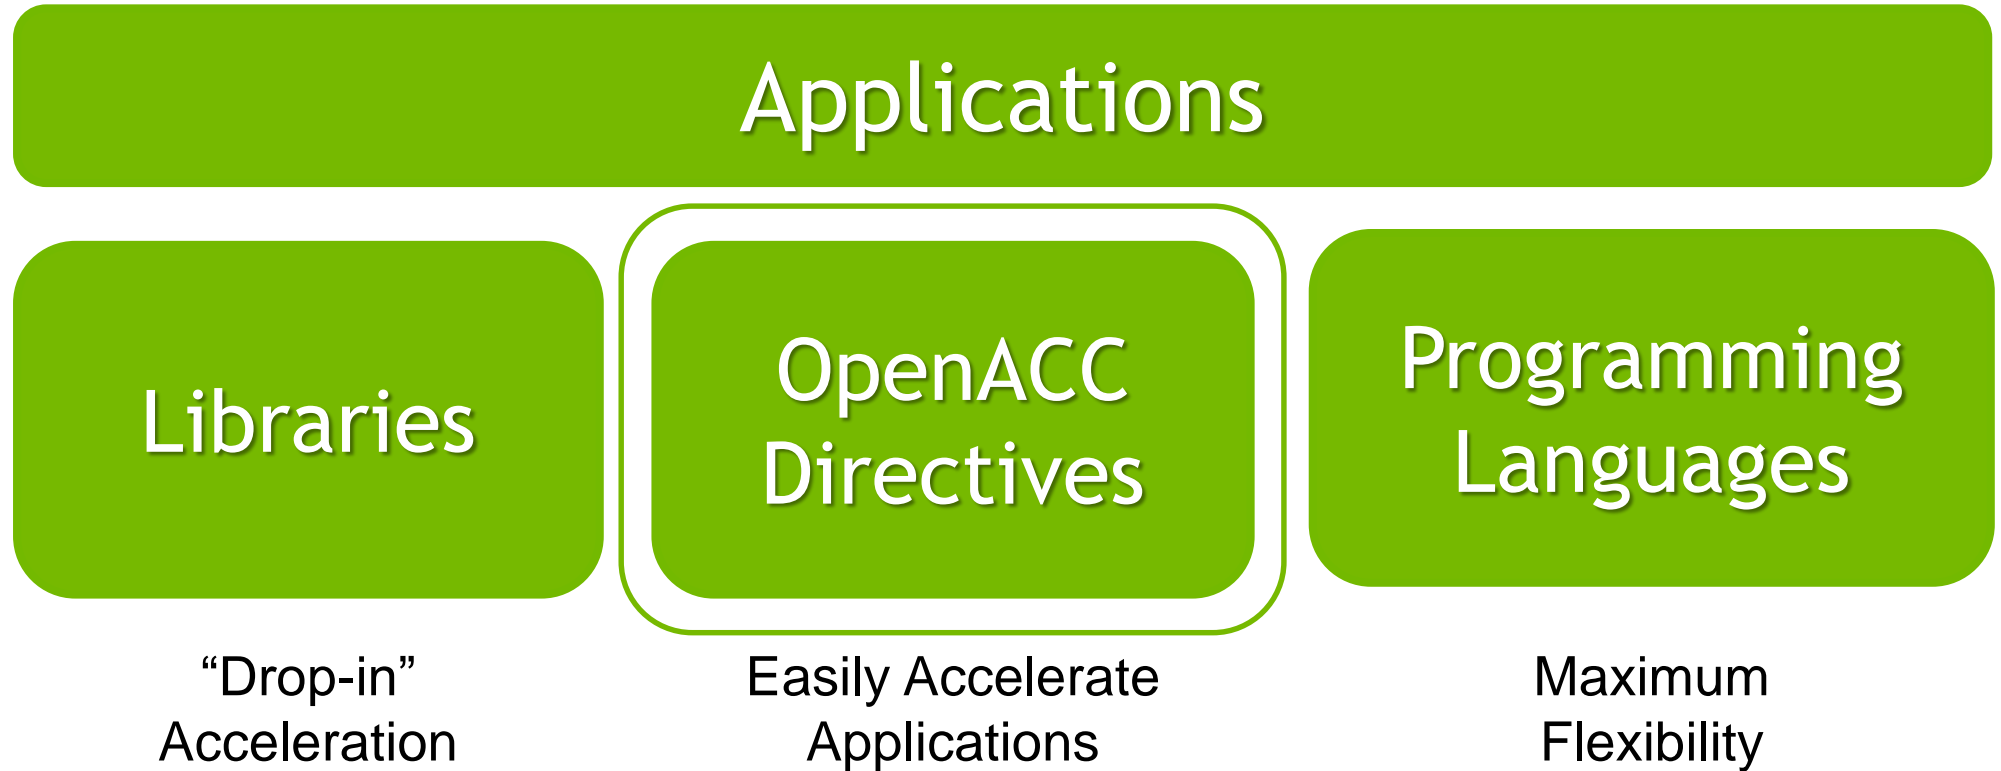
<https://developer.nvidia.com/openacc-advanced-course>

ADVANCED PROFILING OF OPENACC CODE

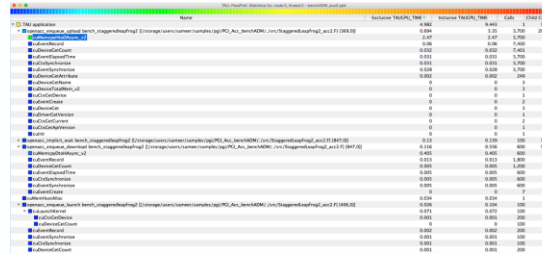
Lecture 1: Ty McKercher, NVIDIA




3 WAYS TO ACCELERATE APPLICATIONS



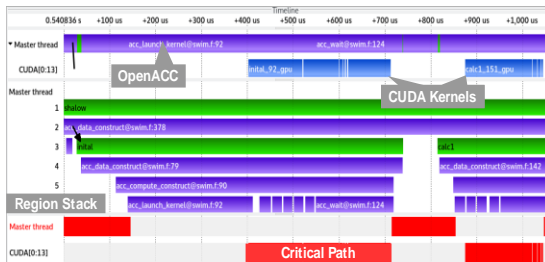
OPENACC PROFILING TOOLS



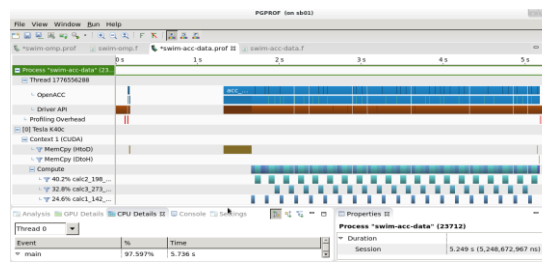
TAU



Vampir



Score-P



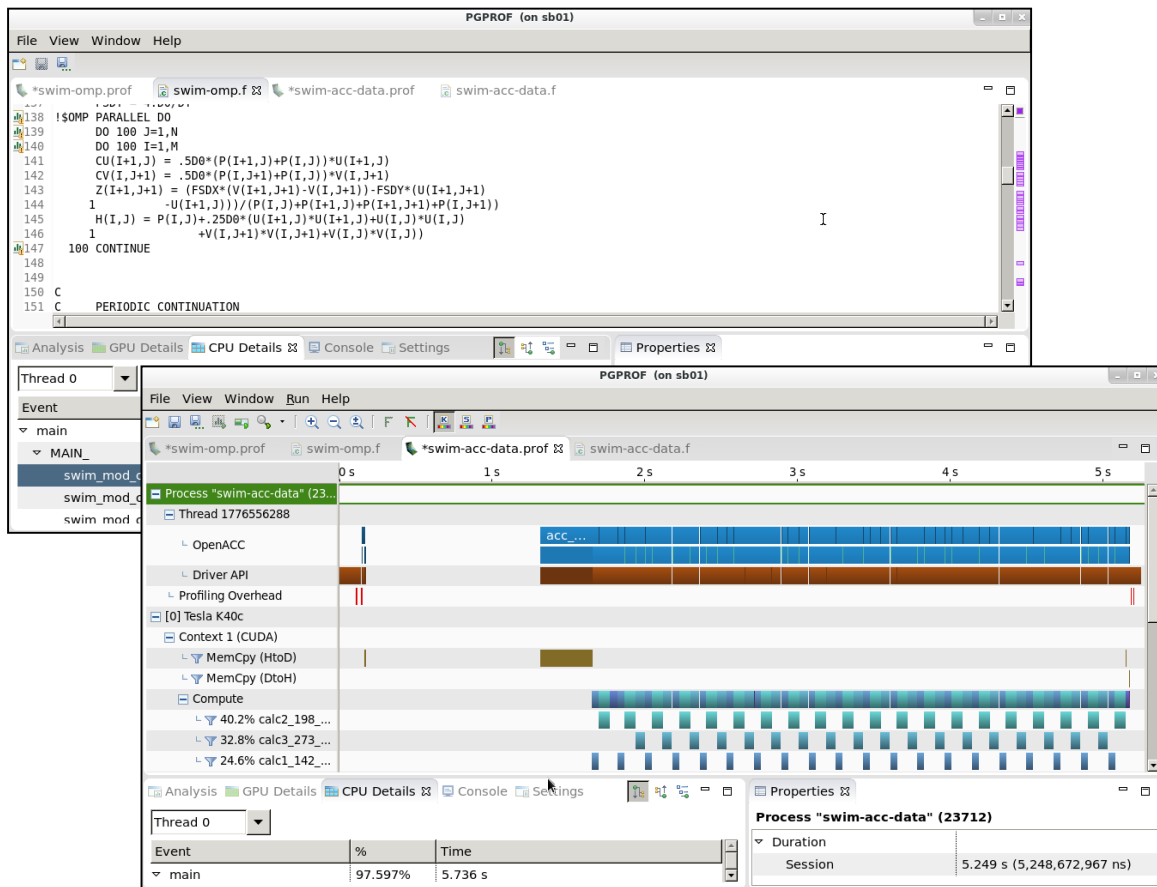
PGPROF

“The OpenACC profiling interface lets us easily measure and investigate implicit operations introduced by the compiler or runtime. **Implementation was straightforward** and basically worked out of the box.”

Dipl.-Ing. Robert Dietrich, T-U Dresden

PGPROF: OPENACC CPU AND GPU PROFILER

Available with the latest OpenACC Toolkit



- For 64-bit multicore processor-based systems with or without accelerators
- Supports thread-level OpenMP profiling
- Supports profiling OpenACC and CUDA Fortran codes on NVIDIA CUDA-enabled GPU accelerators
- Graphical and command-line user interfaces
- Function level (routine) and source code line level profiling
- Comprehensive built-in help facilities

PGPROF (on sb01)

File View Window Help

*swim-omp.prof swim-omp.f *swim-acc-data.prof swim-acc-data.f

```
138 !$OMP PARALLEL DO
139 DO 100 J=1,N
140 DO 100 I=1,M
141 CU(I+1,J) = .5D0*(P(I+1,J)+P(I,J))*U(I+1,J)
142 CV(I,J+1) = .5D0*(P(I,J+1)+P(I,J))*V(I,J+1)
143 Z(I+1,J+1) = (FSDX*(V(I+1,J+1)-V(I,J+1))-FSDY*(U(I+1,J+1)
144 1 -U(I+1,J)))/(P(I,J)+P(I+1,J)+P(I+1,J+1)+P(I,J+1))
145 H(I,J) = P(I,J)+.25D0*(U(I+1,J)*U(I+1,J)+U(I,J)*U(I,J)
146 1 +V(I,J+1)*V(I,J+1)+V(I,J)*V(I,J))
147 100 CONTINUE
148
149
150 C
151 C PERIODIC CONTINUATION
```

Analysis GPU Details CPU Details Console Settings

Thread 0

Event	%	Time
main	100.124%	8.063 s
MAIN_	100.124%	8.063 s
swim_mod_calc3_	32.174%	2.591 s
swim_mod_calc2_	31.677%	2.551 s
swim mod calc1	26.211%	2.111 s

Properties

Select or highlight a single interval to see properties

PGPROF (on sb01)

File View Window Help

*swim-omp.prof swim-omp.f *swim-acc-data.prof swim-acc-data.f

```

138 !$OMP PARALLEL DO
139   Parallel region activated
140   DO 100 I=1,M
141     CU(I+1,J) = .5D0*(P(I+1,J)+P(I,J))*U(I+1,J)
142     CV(I,J+1) = .5D0*(P(I,J+1)+P(I,J))*V(I,J+1)
143     Z(I+1,J+1) = (FSDX*(V(I+1,J+1)-V(I,J+1))-FSDY*(U(I+1,J+1)
144     -U(I+1,J)))/(P(I,J)+P(I+1,J)+P(I+1,J+1)+P(I,J+1))
145     H(I,J) = P(I,J)+.25D0*(U(I+1,J)*U(I+1,J)+U(I,J)*U(I,J)
146     +V(I,J+1)*V(I,J+1)+V(I,J)*V(I,J))
147 100 CONTINUE
148
149
150 C
151 C PERIODIC CONTINUATION

```

Analysis GPU Details CPU Details Console Settings

Thread 0

Event	%	Time
main	100.124%	8.063 s
MAIN_	100.124%	8.063 s
swim_mod_calc3_	32.174%	2.591 s
swim_mod_calc2_	31.677%	2.551 s
swim mod calc1	26.211%	2.111 s

Properties

Select or highlight a single interval to see properties

PGPROF (on sb01)

File View Window Help

*swim-omp.prof swim-omp.f *swim-acc-data.prof swim-acc-data.f

```

138 !$OMP PARALLEL DO
139 DO 100 I=1 N
140 Multiple markers at this line
141   - Intensity = [symbolic], and not printable, try the -Mpf -Mpfo options
142   - Parallel loop activated with static block schedule
143   Z(I+1,J+1) = (FSDX*(V(I+1,J+1)-V(I,J+1))-FSDY*(U(I+1,J+1)
144   1      -U(I+1,J)))/(P(I,J)+P(I+1,J)+P(I+1,J+1)+P(I,J+1))
145   H(I,J) = P(I,J)+.25D0*(U(I+1,J)*U(I+1,J)+U(I,J)*U(I,J)
146   1      +V(I,J+1)*V(I,J+1)+V(I,J)*V(I,J))
147 100 CONTINUE
148
149
150 C
151 C PERIODIC CONTINUATION

```

Analysis GPU Details CPU Details Console Settings

Thread 0

Event	%	Time
main	100.124%	8.063 s
MAIN_	100.124%	8.063 s
swim_mod_calc3_	32.174%	2.591 s
swim_mod_calc2_	31.677%	2.551 s
swim mod calc1	26.211%	2.111 s

Properties

Select or highlight a single interval to see properties

PGPROF (on sb01)

File View Window Help

*swim-omp.prof swim-omp.f *swim-acc-data.prof swim-acc-data.f

```

138 !$OMP PARALLEL DO
139   DO 100 J=1,N
140   DO 100 T=1,M
141
142
143
144
145
146
147   100 CONTINUE
148
149
150 C
151 C   PERIODIC CONTINUATION

```

Multiple markers at this line

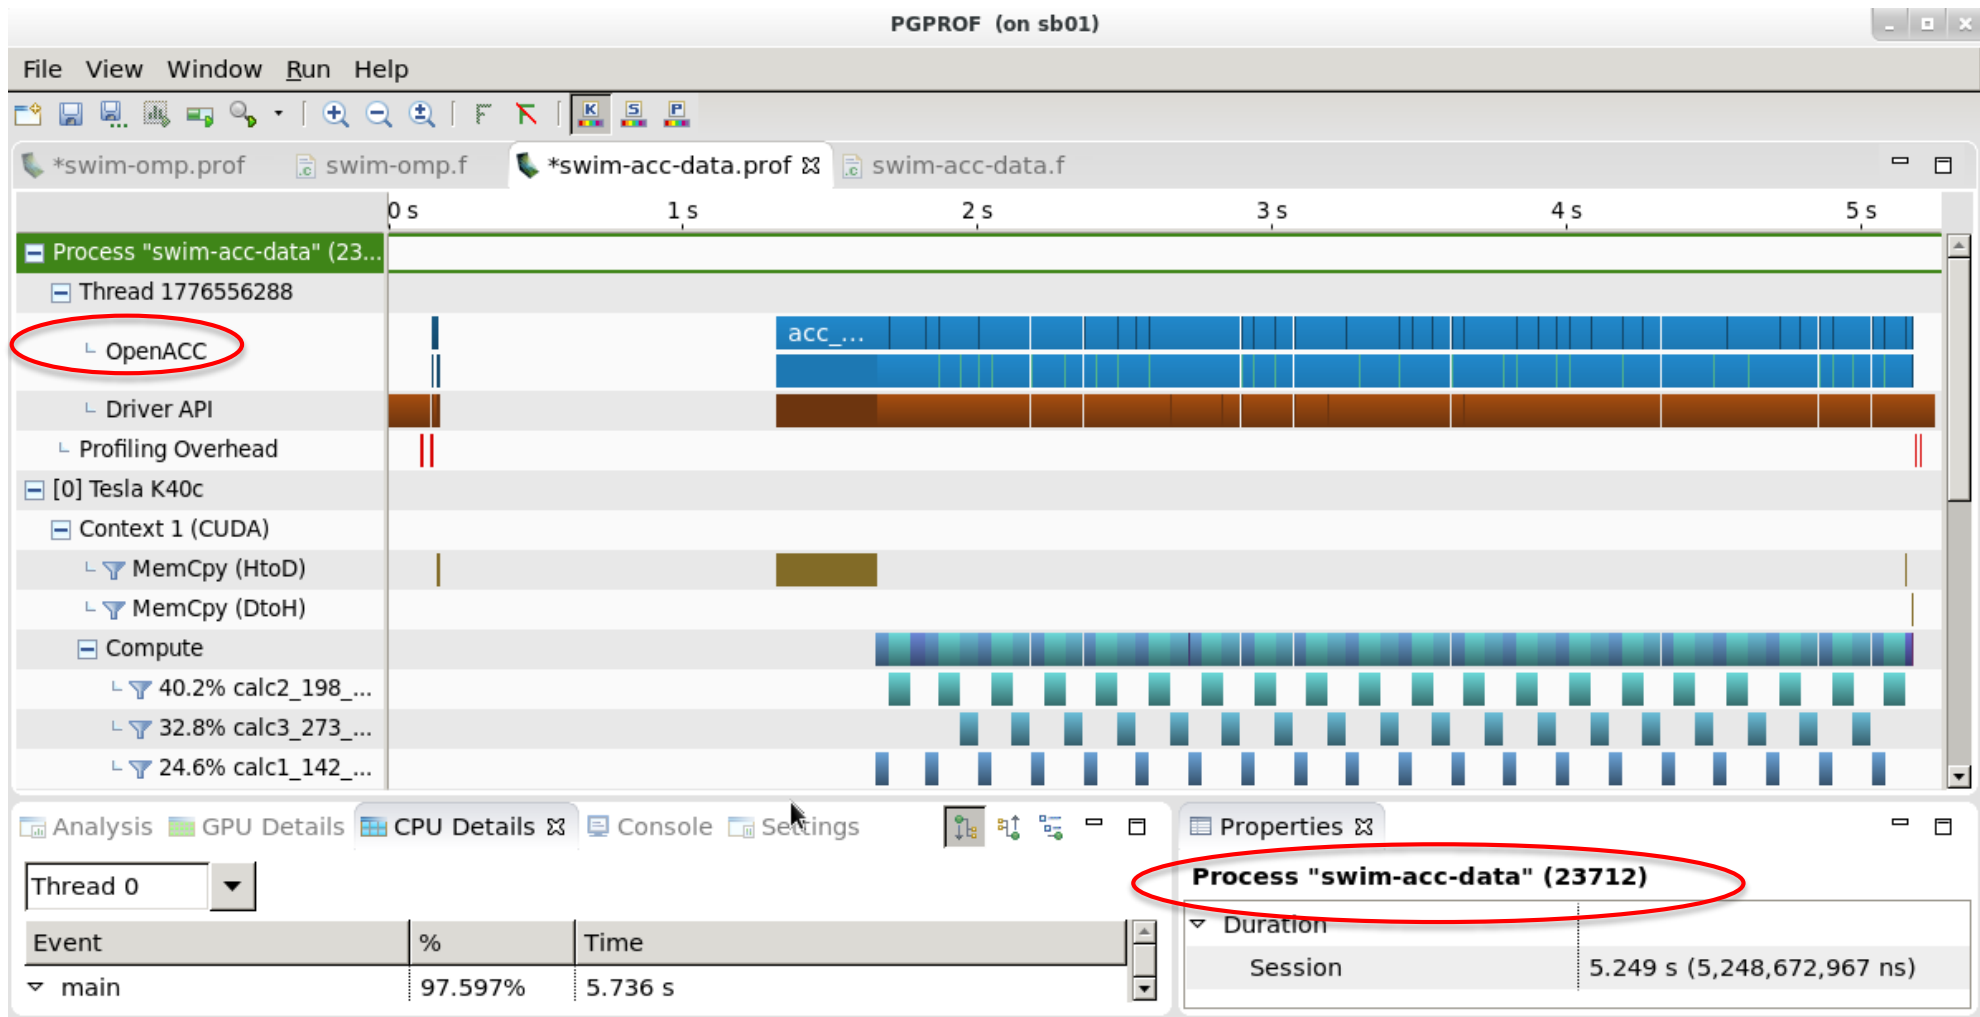
- Generated 6 prefetch instructions for the loop
- Generated vector sse code for the loop
- Generated 5 alternate versions of the loop
- 2 loop-carried redundant expressions removed with 2 operations and 4 arrays
- Intensity = 1.93

Analysis GPU Details CPU Details Console Settings Properties

Thread 0

Event	%	Time
main	100.124%	8.063 s
MAIN_	100.124%	8.063 s
swim_mod_calc3_	32.174%	2.591 s
swim_mod_calc2_	31.677%	2.551 s
swim mod calc1	26.211%	2.111 s

Select or highlight a single interval to see properties



More on Visual Profiler Office Hour on May 26th

PGPROF Quick Start Guide: <http://www.pgroup.com/resources/pgprof-quickstart.htm>

PGPROF User's Guide*: <http://www.pgroup.com/doc/pgprofug.pdf>

Agenda

Profiling application - Getting Started

Acceleration with Managed Memory

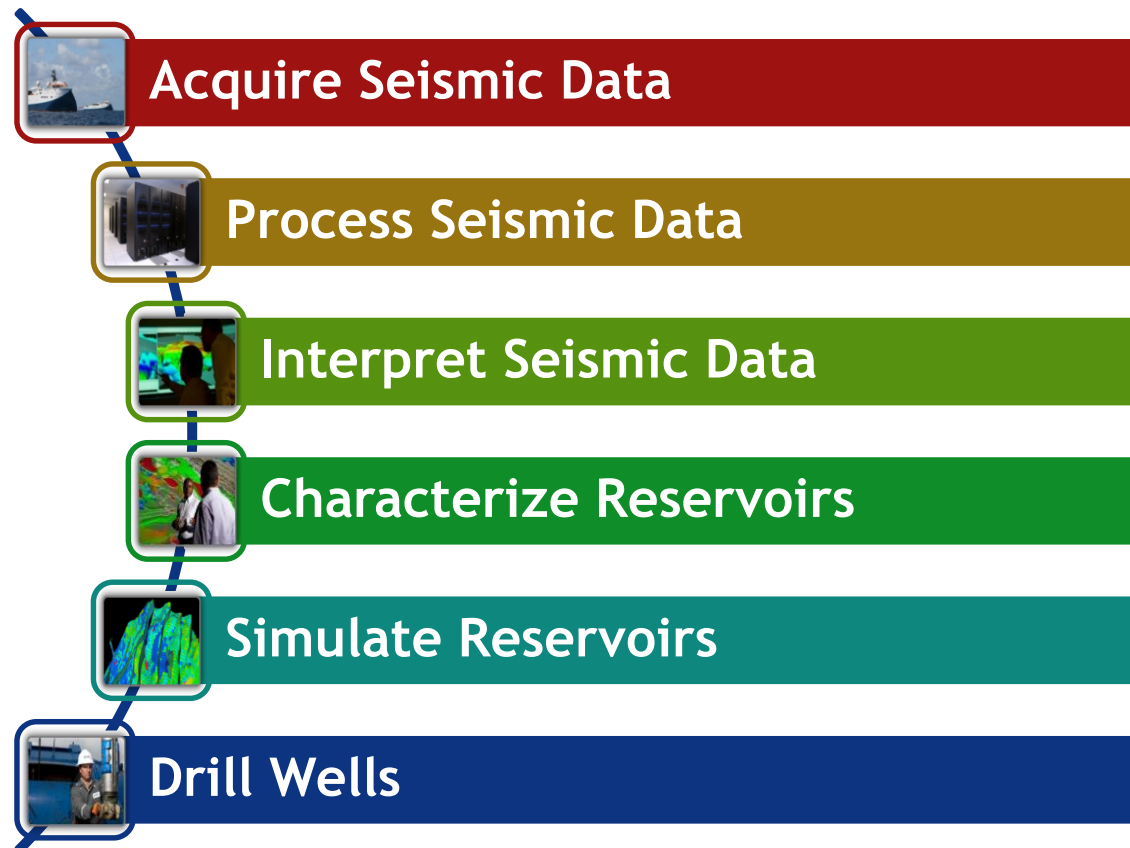
Optimization with Data Directives

Multicore comparison

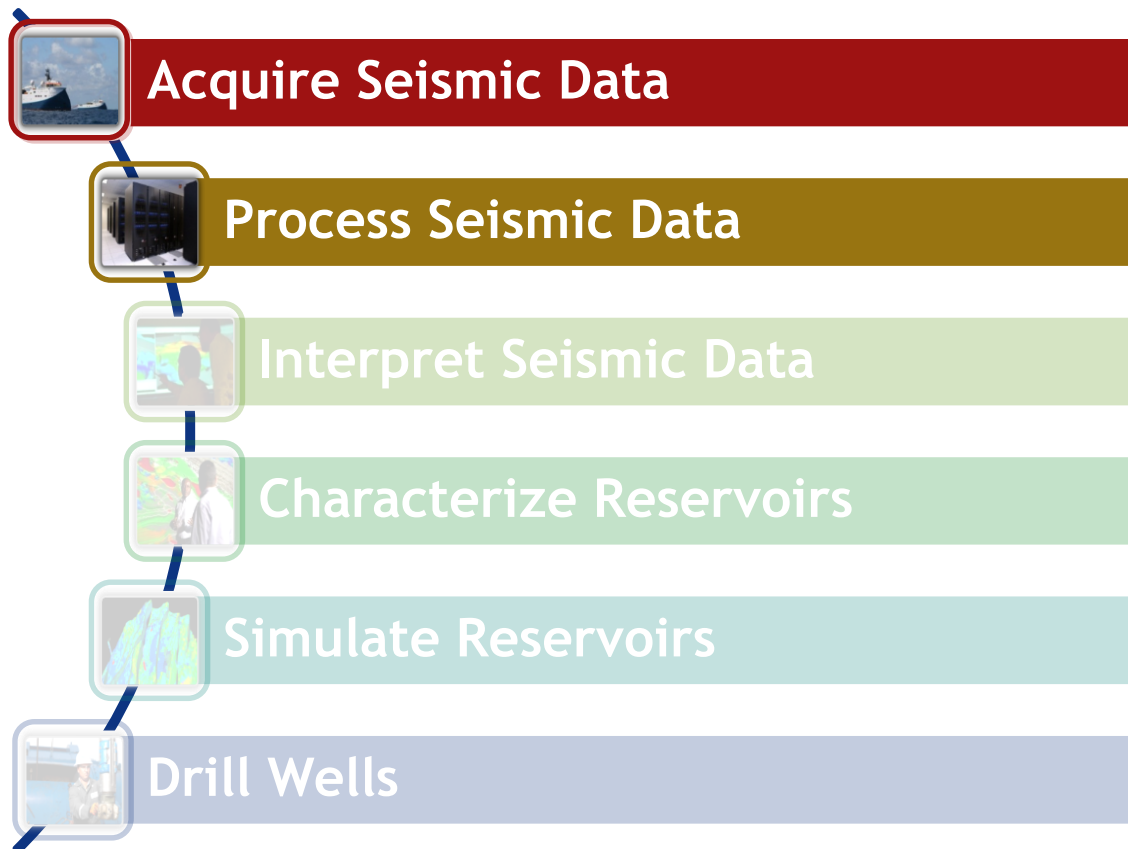
Seismic Unix Configuration

Profiling Application - Getting Started


EXPLORATION & PRODUCTION WORKFLOW



EXPLORATION & PRODUCTION WORKFLOW




ACQUIRE SEISMIC DATA



REAL ROCK vs SEISMIC REFLECTION


Imaging, signal processing, filtering, ray tracing



USE KNOWN MODEL

Verify accuracy of imaging algorithms

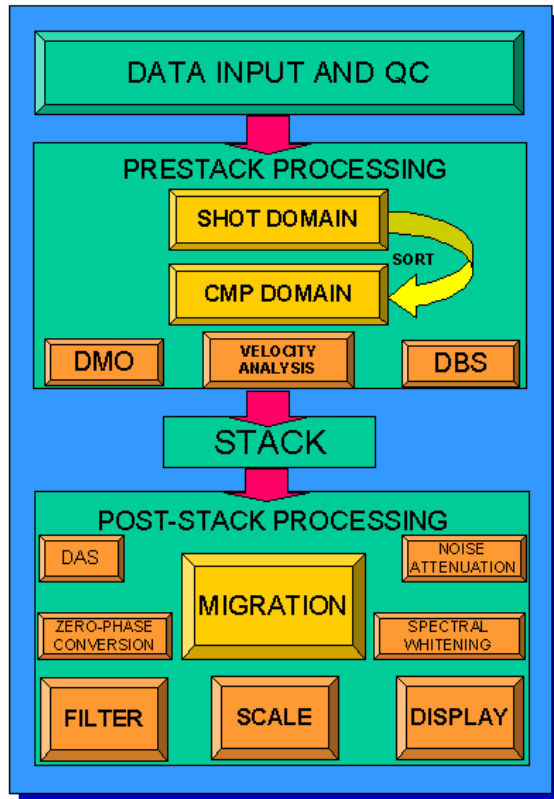
- ▶ Industry standard dataset
- ▶ Compare against known truth
- ▶ Refine seismic algorithms
- ▶ Improve confidence in drilling decisions



PROCESS SEISMIC DATA

Open Source Seismic*Unix Package

Typical Seismic Processing Flow



- ▶ Iterative process
 - ▶ Can take several months
- ▶ Migration most time consuming
 - ▶ Collapse diffractions
 - ▶ Correctly position dip events
- ▶ Trusted tool: Kirchhoff Migration


PSEUDO-CODE FOR KIRCHHOFF MIGRATION

sukdmig2d

- Update residual travel times
- Loop over traces
 - Combine travel times (sum2)
 - Migrate trace (mig2d)
 - Low-pass filter trace
 - Compute amplitudes
 - Interpolate along lateral
 - Interpolate along depth
- Write output

```
main() {  
    ...  
    resit();  
  
    do {  
        ...  
        sum2(...,ttab, tt);  
        sum2(...,tsum, tsum);  
  
        mig2d(..., mig, ..., mig1, ...);  
    } while (fget(tr) && jtr<ntr);  
  
    fput(tr);  
  
}
```

HOW YOU ACCELERATE CODE WITH GPUS



ASSESS BASELINE CPU PROFILE

Use pgprof sukdmig2d

The screenshot shows the PGPROF application interface. The main window displays the source code of 'sukdmig2d.c', which includes a large loop for processing data. The bottom panel contains a console window with the following output:

```
<terminated> /cm/extra/apps/PGI/16.4/linux86-64/2016/cuda/7.5-pgprof/bin/nvprof
noff=1 off0=0 doff=99999
v0=1500 dvz=0
aperx=4612.5 offmax=3000 angmax=60
ntr=100000 mtr=1 ls=1 npv=0
absoff=0 limoff=0
=====
input traveltime tables
compute traveltime residual ...
start migration ...

fs=0 es=9200 offmax=3000
migrated 23040 traces in total

output done
```

Function	Percent Runtime
mig2d	79%
sum2	8.5 %
resit	< 1%

RESULTS

SUKDMIG2D	Configuration	Model Size	Cores	Elapsed Time (s)	Speed up
CPU Only (Baseline)	2x E5-2698 v3 2.30GHz	2301 x 751	1	218	1.00

Acceleration with Managed Memory

PARALLELIZE

Parallel Directives

`#pragma`

Parallelize `for` loops

Vectorize

Compiler vectorizes inner loops

```
mig2d:
    #pragma acc parallel for
    for (ix=nxtf; ix<=nxte; ++ix) {
    . . .
        #pragma acc loop
        for (iz=izt0; iz<nzt; ++iz) {
        . . .
```

PARALLELIZE

Parallel Directives

`restrict` on pointers!

limits aliasing

www.wikipedia.org/wiki/Restrict

`#pragma`

Parallelize outer `for` loops

Compiler parallelizes inner loop

```
void sum2(int nx, int nz, float a1, float a2,
          float ** restrict t1, float ** restrict t2, float **
          restrict t)
{
    int ix, iz;

    #pragma acc parallel for
    for(ix=0; ix < nx; ++ix)
    {
        for(iz=0; iz < nz; ++iz)
            t[ix][iz] = a1*t1[ix][iz]+a2*t2[ix][iz];
    }
}
```

PARALLELIZE

Resolve Errors!

Parallel Directives

`#pragma`

Parallelize `for` outer loop


Parallelize inner loops

Resolve loop carried depend

Add acc loop directive


```
Resit (managed):  
537, Accelerator kernel generated  
Generating Tesla code  
538, #pragma acc loop gang /* blockIdx.x */  
553, #pragma acc loop vector(128) /* threadIdx.x */  
540, Loop carried dependence of t->-> prevents parallelization  
Loop carried backward dependence of t->-> prevents vector
```

```
Resit:  
#pragma acc parallel for  
for (ix=0; ix<nx; ++ix)  
{  
    #pragma acc loop  
    for (is=0; is<ns; ++is)  
    {  
        . . .  
        #pragma acc loop  
        for (iz=0; iz<nz; ++iz)  
            t[ix][iz] -= sr0*tb[jr][iz]+sr*tb[jr+1][iz];
```




UNIFIED MEMORY IMPROVES PRODUCTIVITY

Previous Developer View



Developer View With Unified Memory



OPENACC AND UNIFIED MEMORY

All heap allocations are in managed memory (Unified Memory Heap)

Pointers can be used on GPU and CPU

Enabled with compiler switch `-ta=tesla:managed`,...

More Info at „OpenACC and CUDA Unified Memory”, by Michael Wolfe, PGI Compiler Engineer: <https://www.pgroup.com/lit/articles/insider/v6n2a4.htm>

OPENACC AND UNIFIED MEMORY

Advantages

No need for any data clauses

No need to fully understand application data flow and allocation logic

Incremental profiler-driven acceleration possible

Outlook to GPU programming on Pascal

RESULTS

SUKDMIG2D	Configuration	Model Size	Cores	Elapsed Time (s)	Speed up
CPU Only (Baseline)	2x E5-2698 v3 2.30GHz	2301 x 751	1	218	1.00
OpenACC (Managed)	1x K40 GPU	2301 x 751	2880	46	4.70

RESULTS

SUKDMIG2D	Configuration	Model Size	Cores	Elapsed Time (s)	Speed up
CPU Only (Baseline)	2x E5-2698 v3 2.30GHz	2301 x 751	1	218	1.00
OpenACC (Managed)	1x K40 GPU	2301 x 751	2880	46	4.70

Now optimize using the Verbose output from compiler!

Optimization with Data Directives

OPTIMIZATION

Compile

```
pgcc -acc \  
-ta=tesla:managed
```

Profile !

```
pgprof <managed binary>
```

==55246== Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
42.82%	4.03645s	23040	175.19us	121.12us	196.38us	mig2d_787_gpu
28.79%	2.71389s	23040	117.79us	80.800us	135.68us	mig2d_726_gpu
27.35%	2.57762s	69120	37.291us	33.248us	42.240us	sum2_571_gpu
1.00%	93.936ms	23040	4.0770us	3.2000us	12.992us	[CUDA memcpy HtoD]
0.04%	3.4627ms	1	3.4627ms	3.4627ms	3.4627ms	resit_537_gpu
0.00%	126.14us	1	126.14us	126.14us	126.14us	timeb_592_gpu

==55246== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
30.16%	11.5982s	230423	50.334us	118ns	3.9101ms	cuMemFree
29.21%	11.2327s	230429	48.746us	10.132us	12.821ms	cuMemAllocManaged
27.15%	10.4430s	253444	41.204us	1.0420us	3.4680ms	cuStreamSynchronize
10.42%	4.00751s	115202	34.786us	5.4290us	99.805ms	cuLaunchKernel
1.13%	433.50ms	1428513	303ns	141ns	429.42us	cuPointerGetAttrib
0.81%	310.55ms	1	310.55ms	310.55ms	310.55ms	cuDevicePrimary...
0.71%	273.10ms	23040	11.853us	7.3210us	409.13us	cuMemcpyHtoDAsync
0.33%	125.36ms	1	125.36ms	125.36ms	125.36ms	cuDevicePrimary...
0.06%	24.165ms	1	24.165ms	24.165ms	24.165ms	cuMemHostAlloc...
0.02%	9.5668ms	1	9.5668ms	9.5668ms	9.5668ms	cuMemFreeHost
0.00%	534.34us	1	534.34us	534.34us	534.34us	cuMemAllocHost
0.00%	461.71us	1	461.71us	461.71us	461.71us	cuModuleLoad..
0.00%	363.83us	2	181.91us	180.02us	183.81us	cuMemAlloc

OPTIMIZATION

Managed Compile

Verbose output
Guided enhancements
Targeted changes

Common Optimizations

Data Movement
Copy, copyin, copyout
Create, delete
Update

Loop Collapse

```
main:
  453, Generating update host(mig[:noff][:nxo][:nzo])
  455, Generating update host(mig1[:noff][:1][:1])
  459, Generating update host(mig1[:noff][:nxo][:nzo])

resit:
  539, Generating copyin(ttab[:ns],tb[:][:nz])

sum2:
  571, Generating copyin(t2[:nx][:nz],t1[:nx][:nz])
      Generating copyout(t[:nx][:nz])

mig2d:
  721, Generating copy(ampt1[nxtf:nxte-nxtf+1][:])
      Generating copyin(cssum[nxtf:nxte-nxtf+1][:],tvsum[nxtf:nxte-nxtf+1][
      Generating copy(tmt[nxtf:nxte-nxtf+1][:],ampti[nxtf:nxte-nxtf+1][:])
      Generating copyin(pb[:][:])
      Generating copy(ampt[nxtf:nxte-nxtf+1][:])
      Generating copyin(cs0b[:][:],angb[:][:])
      Generating copy(zpt[nxtf:nxte-nxtf+1])
  782, Generating copy(mig1[nxf:nxe-nxf+1][:])
      Generating copyin(ampt1[:][:], tb[:][:], tsum[:][:], ampt[:][:], ...
      Generating copy(mig[nxf:nxe-nxf+1][:])
      Generating copyin(zpt[:])
```

OPTIMIZATION

Data Movement

Analyze data flow in application
Explicitly use data directives

Move data directive to *main*

Create only when possible

Copyin move data to GPU

Update to move data to host

```
main:
#pragma acc enter data create(tb,pb,cs0b,ang0)
#pragma acc enter data create(tt,tsum)
#pragma acc enter data copyin(mig, ttab)

#pragma acc enter data create(tvsum,csum )
#pragma acc enter data copyin(cs, tv)
#pragma acc enter data copyin(mig1)

After processing:

#pragma acc update host(mig)
#pragma acc update host(mig1)
```

OPTIMIZATION

Data Movement

Explicitly use present for data already on GPU!

Collapse

Increase the threads $nx*nz$

Present

Data is already on the GPU
Prevent data movement

```
sum2: (managed)
571, Generating copyin(t2[:nx][:nz],t1[:nx][:nz])
Generating copyout(t[:nx][:nz])
```

```
void sum2(int nx, int nz, float a1, float a2,
         float ** restrict t1, float ** restrict t2, float **
         restrict t)
{
    int ix, iz;

    #pragma acc parallel for collapse(2) present(t1,t2,t)
    for(ix=0; ix < nx; ++ix)
    {
        for(iz=0; iz < nz; ++iz)
            t[ix][iz] = a1*t1[ix][iz]+a2*t2[ix][iz];
    }
}
```

OPTIMIZATION

Data Movement

Move large data transfers to main i.e. mig, mig1

Minimize Copyin, Copyout

Maximize Create, Present
Prevents data transfers

Use Copyin, Copyout, Copy
only when data changes!

Delete happens when
leaving scope

```
mig2d:
721, Generating copy(ampt1[nxtf:nxte-nxtf+1][:])
Generating copyin(cssum[nxtf:nxte-nxtf+1][:],tvsum[...
Generating copy(tmt[nxtf:nxte-nxtf+1][:],ampti[...
Generating copyin(pb[:][:])
Generating copy(ampt[nxtf:nxte-nxtf+1][:])
Generating copyin(cs0b[:][:],angb[:][:])
Generating copy(zpt[nxtf:nxte-nxtf+1])
782, Generating copy(mig1[nxf:nxe-nxf+1][:])
Generating copyin(ampt1[:][:], tb[:][:], tsum[:][:], ...
Generating copy(mig[nxf:nxe-nxf+1][:])
```

```
void mig2d(float * restrict trace, int nt, float ft,...)
{
...
#pragma acc data
copyin(trace[0:nz],trf[0:nt+2*mtmax]) \
present(mig, mig1, tb,tsum,tvsum,cssum,pb,... \
create(tmt[0:nxt][0:nzt], ampt[0:nxt][0:nzt],...
{
```


OPTIMIZATION

Data Movement

Use present for data already on GPU!

Collapse

Increase the threads $nx*ns$

Present

Data is already on the GPU
Prevent data movement

```
Resit: (managed)
539, Generating copyin(ttab[:ns],tb[:][:nz])
```

```
resit:
...
#pragma acc parallel for collapse(2) present(tb, ttab)
for (ix=0; ix<nx; ++ix)
{
    for (is=0; is<ns; ++is)
    {
...
        #pragma acc loop
        for (iz=0; iz<nz; ++iz)
            t[ix][iz] -= sr0*tb[jr][iz]+sr*tb[jr+1][iz];
    }
}
```

OPTIMIZATION

Data Movement

mig, mig1 data large

Move to main

Copyin at start

Mark as present

Copyout for snapshots

Minimize Copyin, Copyout

Use create

Prevents copy in/out

Delete happens when leaving scope

```
void mig2d(float * restrict trace, int nt, float
ft,...)
{
...
#pragma acc data
    copyin(trace[0:nz],trf[0:nt+2*mtmax]) \
    present(mig, mig1, tb,tsum,tvsum,cssum,pb,... \
    create(tmt[0:nxt][0:nzt], ampt[0:nxt][0:nzt],...
    {
...
        #pragma acc parallel for
        for (ix=nxtf; ix <= nxte; ++ix) {
...
            #pragma acc loop
            for (iz=izt0; iz < nzt; ++iz) {
...

```

OPTIMIZATION

Compile

```
pgcc -acc -ta=tesla
```

Profile

```
pgprof <tesla binary>
```

mig2d and sum2 about the same.

- ✓ cuAllocManaged (11s) removed.
- ✓ cuMemFree (11.5s) reduced to milliseconds.

==2242== Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
41.54%	3.95071s	23040	171.47us	118.88us	192.61us	mig2d_787_gpu
27.91%	2.65415s	23040	115.20us	78.241us	133.09us	mig2d_726_gpu
26.27%	2.49826s	69120	36.143us	32.768us	40.416us	sum2_569_gpu
2.88%	274.19ms	69132	3.9660us	3.5520us	13.120us	__pgi_uacc_cuda_fill_32_gpu
1.35%	128.68ms	46088	2.7920us	2.4960us	1.6815ms	[CUDA memcpy HtoD]
0.04%	3.4187ms	1	3.4187ms	3.4187ms	3.4187ms	resit_535_gpu
0.00%	226.15us	2	113.07us	2.4640us	223.68us	[CUDA memcpy DtoH]
0.00%	123.43us	1	123.43us	123.43us	123.43us	timeb_592_gpu

==2242== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
85.89%	9.71880s	138246	70.300us	1.8870us	3.4228ms	cuStreamSynchronize
7.69%	869.62ms	184334	4.7170us	3.4420us	452.72us	cuLaunchKernel
2.94%	333.00ms	1	333.00ms	333.00ms	333.00ms	cuDevicePrimaryCtxRetain
1.75%	197.59ms	46088	4.2870us	2.8370us	426.78us	cuMemcpyHtoDAsync
1.15%	130.58ms	1	130.58ms	130.58ms	130.58ms	cuDevicePrimaryCtxRelease
0.25%	28.337ms	1	28.337ms	28.337ms	28.337ms	cuMemHostAlloc
0.20%	23.059ms	46084	500ns	260ns	11.292us	cuPointerGetAttributes
0.09%	10.027ms	1	10.027ms	10.027ms	10.027ms	cuMemFreeHost
0.03%	2.9512ms	31	95.199us	2.9220us	300.63us	cuMemAlloc
0.01%	806.38us	2	403.19us	188.55us	617.83us	cuModuleLoadData

No longer compiling with :managed

RESULTS

SUKDMIG2D	Configuration	Model Size	Cores	Elapsed Time (s)	Speed up
CPU Only (Baseline)	2x E5-2698 v3 2.30GHz	2301 x 751	1	218	1.00
OpenACC (GPU Managed)	1x K40 GPU	2301 x 751	2880	46	4.70
OpenACC (GPU Native)	1x K40 GPU	2301 x 751	2880	12	15.60

Multicore Comparison

OPTIMIZATION

How about Multi-Core / OMP / pthread?

Done!

Re-Compile

```
pgcc -acc -ta=multicore
```

Profile !

```
pgprof
--cpu-profiling on \
--cpu-profiling-scope function \
--cpu-profiling-mode top-down \
<app> <args>
```


```
===== CPU profiling result (top down):
72.91% main
| 69.84% mig2d
| | 43.19% __pgi_acc_barrier
| | | 43.19% _mp_barrier_tw
| | | 0.02% _mp_pcpu_get_team_lcpu
| | | 0.02% _mp_pcpu_struct
| | | 0.01% __tls_get_addr
| | 0.12% malloc@@GLIBC_2.2.5
| 2.88% sum2
| | 2.79% __pgi_acc_barrier
| | | 2.79% _mp_barrier_tw
| | 0.00% .ACCENTER
| | 0.00% _mp_barrierr
| 0.10% __fsd_cos_vex
| 0.05% __pgi_acc_pexit
| | 0.05% _mp_cpexit
| | 0.05% _mp_barrierw
22.18% _mp_slave
| 22.18% _mp_cslave
| 22.18% _mp_barrier_tw
| 0.02% _mp_pcpu_yield
| 0.02% sched_yield
4.77% __fsd_cos_vex
0.09% filt
-- more --
```

RESULTS


SUKDMIG2D	Configuration	Model Size	Cores	Elapsed Time (s)	Speed up
CPU Only (Baseline)	2x E5-2698 v3 2.30GHz	2301 x 751	1	218	1.00
OpenACC CPU (Multicore)	2x E5-2698 v3 2.30GHz	2301 x 751	16	29	7.50
OpenACC GPU (Managed)	1x K40 GPU	2301 x 751	2880	46	4.70
OpenACC GPU (Native)	1x K40 GPU	2301 x 751	2880	12	15.60

DEPLOY


How do the results compare?



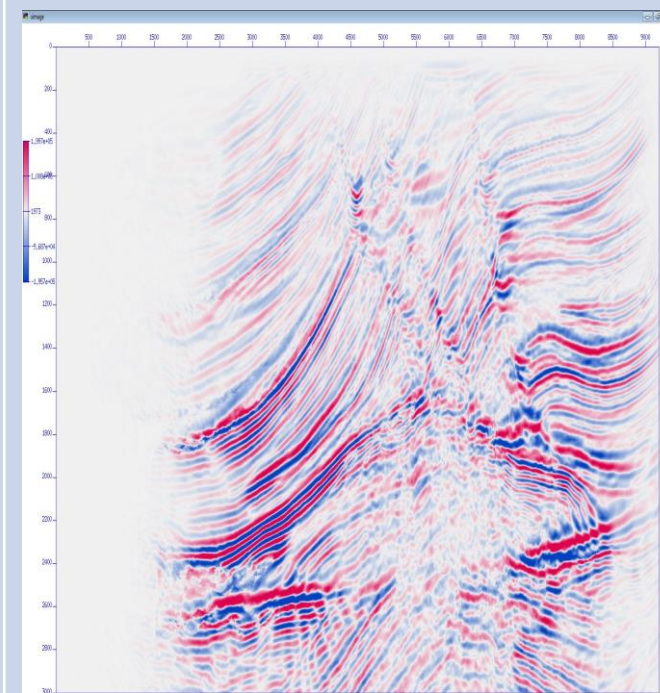
CPU Only (Baseline)



OpenACC Multicore



OpenACC (GPU)



Homework

QWIKLABS: GETTING ACCESS

1. Go to <https://developer.nvidia.com/qwiklabs-signup>
2. Register with OpenACC promo code to get free access
3. Receive a confirmation email with access instructions

Questions?

Email to openacc@nvidia.com

ACCESS TO HOMEWORK

Qwiklab:

Profile-driven approach to accelerate Seismic application with OpenACC

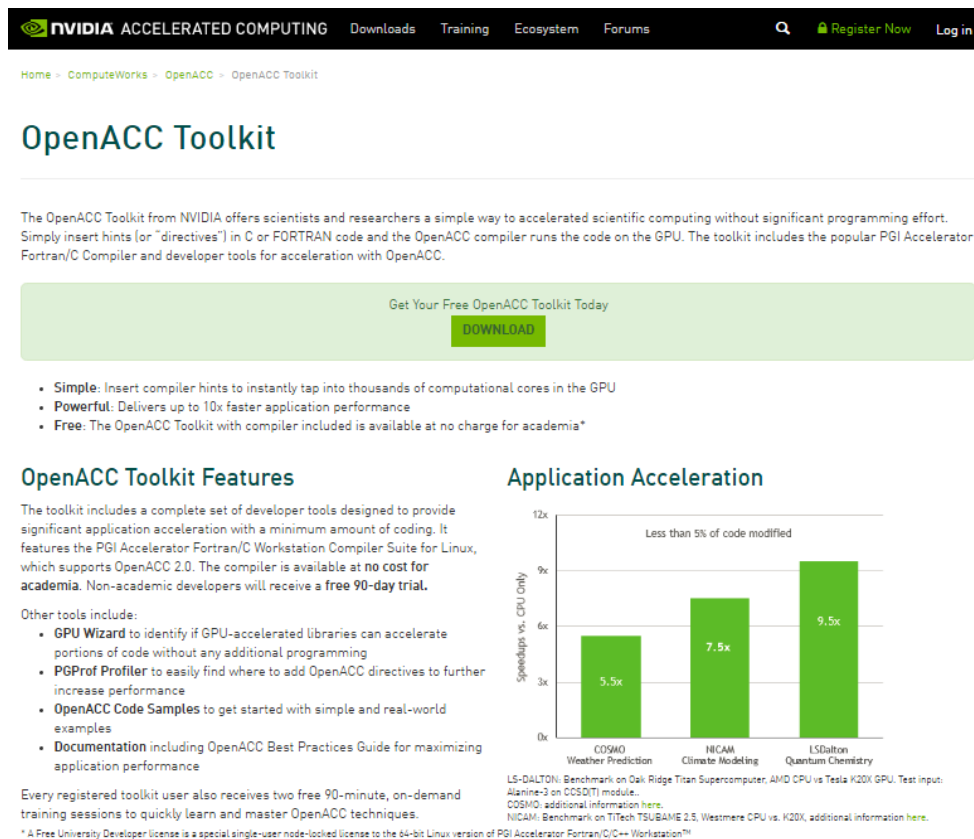
Link: <http://bit.ly/oaccnvlab6>

Requirements: OpenACC Compiler and CUDA-aware MPI

Link to the source code on github thorough the qwiklab if you want to try it on your machine

INSTALL THE OPENACC TOOLKIT (OPTIONAL)

- ▶ Go to developer.nvidia.com/openacc-toolkit
- ▶ Register for the OpenACC Toolkit
- ▶ Install on your personal machine (Linux Only)
- ▶ Free workstation license for academia/90 day free trial for the rest



The screenshot shows the NVIDIA Accelerated Computing website for the OpenACC Toolkit. The header includes the NVIDIA logo and navigation links: Downloads, Training, Ecosystem, Forums, a search icon, and links to Register Now and Log in. The breadcrumb trail is Home > ComputeWorks > OpenACC > OpenACC Toolkit. The main heading is "OpenACC Toolkit". Below it, a paragraph describes the toolkit as a simple way to accelerate scientific computing without significant programming effort, mentioning the PGI Accelerator Fortran/C Compiler and developer tools. A green button labeled "DOWNLOAD" is prominent. Below the button, three bullet points highlight the toolkit's features: Simple (inserting hints for GPU acceleration), Powerful (up to 10x performance), and Free (available at no charge for academia). The "OpenACC Toolkit Features" section describes the complete set of developer tools, including the PGI Accelerator Fortran/C Workstation Compiler Suite for Linux, which supports OpenACC 2.0. It also lists other tools: GPU Wizard, PGProf Profiler, OpenACC Code Samples, and documentation. To the right, the "Application Acceleration" section features a bar chart showing speedups for COSMO, NICAM, and LS-DALTON benchmarks. A note states "Less than 5% of code modified".

OpenACC Toolkit Features

The toolkit includes a complete set of developer tools designed to provide significant application acceleration with a minimum amount of coding. It features the PGI Accelerator Fortran/C Workstation Compiler Suite for Linux, which supports OpenACC 2.0. The compiler is available at **no cost for academia**. Non-academic developers will receive a **free 90-day trial**.

Other tools include:

- **GPU Wizard** to identify if GPU-accelerated libraries can accelerate portions of code without any additional programming
- **PGProf Profiler** to easily find where to add OpenACC directives to further increase performance
- **OpenACC Code Samples** to get started with simple and real-world examples
- **Documentation** including OpenACC Best Practices Guide for maximizing application performance

Every registered toolkit user also receives two free 90-minute, on-demand training sessions to quickly learn and master OpenACC techniques.

Application Acceleration

Less than 5% of code modified

Application	Speedups vs. CPU Only
COSMO Weather Prediction	5.5x
NICAM Climate Modeling	7.5x
LS-DALTON Quantum Chemistry	9.5x

LS-DALTON: Benchmark on Oak Ridge Titan Supercomputer, AMD CPU vs Tesla K20X GPU. Test input: Algaline-3 on CCSD(T) module.
COSMO: additional information [here](#).
NICAM: Benchmark on TITech TSUBAME 2.5, Westmere CPU vs. K20X, additional information [here](#).

* A Free University Developer license is a special single-user node-locked license to the 64-bit Linux version of PGI Accelerator Fortran/C/C++ Workstation™.



On Your Own ... Local Seismic Unix Setup

SETUP SEISMIC UNIX

Center for Wave Phenomena

Download Seismic Unix

ftp://ftp.cwp.mines.edu/pub/cwpcodes/cwp_su_all_43R8.tgz

Unpack to ~/cwp

Set environment variables

CWPROOT=~/cwp

PATH=~/cwp/bin:\$PATH

Edit Makefile.config, build

Use PGI compilers (CC=pgcc, FC=pgfortran)

OPTC=-g, FFLAGS=\$(FOPTS)

SETUP SEISMIC UNIX

Marmousi Datasets

Download Marmousi data, velocity, and density files

<http://www.trip.caam.rice.edu/downloads/ieee.tar.gz>

Convert SEGY format to SU format

```
#!/bin/bash
segypread tape=data.segy      conv=0 endian=0 > data.su
segypread tape=velocity.segy conv=0 endian=0 > velocity.su

suflip flip=0 < velocity.su > velocity1.su
sustrip < velocity1.su > velocity.h@ ftn=0

suwind < data.su > data1.su tmax=2.9
```

SETUP SEISMIC UNIX

Smooth, build ray trace model, migrate

```
#!/bin/bash
```

```
nz = 751
```

```
nx = 2301
```

```
dz = 4
```

```
dx = 4
```

```
nt = 750
```

```
ntr= 96
```

```
dt = 4000
```

```
ifile = data1.su
```

```
ofile = datamig.su
```

```
tfile = tfile
```

```
vfile = velocity.h@
```

```
#smoothing
```

```
time smooth2 < $vfile n1=$nz n2=$nx r1=20 r2=20
```

```
>smoothvel
```

```
#raytrace
```

```
time rayt2d < smoothvel dt=0.004 nt=751 dz=$dz nz=$nz
```

```
dx=$dx nx=$nx fxo=0 dxo=25 nxo=369 fxs=0 dxs=100 nxs=93
```

```
>$tfile
```

```
#migrate (Example)
```

```
sukdmig2d infile=$ifile datain=$ifile outfile=$ofile
```

```
dataout=$ofile ttfile=$tfile fzt=0 dzt=4 nzt=751 fxt=0
```

```
dxt=25 nxt=369 fs=0 ns=93 ds=100 nzo=751 dzo=4 dxm=25
```

```
mtr=1
```


WHERE TO FIND HELP

- OpenACC Course Recordings - <https://developer.nvidia.com/openacc-courses>
- PGI Website - <http://www.pgroup.com/resources>
- OpenACC on StackOverflow - <http://stackoverflow.com/questions/tagged/openacc>
- OpenACC Toolkit - <http://developer.nvidia.com/openacc-toolkit>
- Parallel Forall Blog - <http://devblogs.nvidia.com/parallelforall/>
- GPU Technology Conference - <http://www.gputechconf.com/>
- OpenACC Website - <http://openacc.org/>

Questions? Email openacc@nvidia.com

Course Syllabus

May 19: Advanced Profiling of OpenACC Code

May 26: Office Hours <- **Visual Profiler**

June 2: Advanced multi-GPU Programming with MPI and OpenACC

June 9: Office Hours

Recordings:

<https://developer.nvidia.com/openacc-advanced-course>

Questions? Email openacc@nvidia.com